

Performance Management

Fundamentals of IdentityIQ
Implementation

IdentityIQ 7.0

Overview

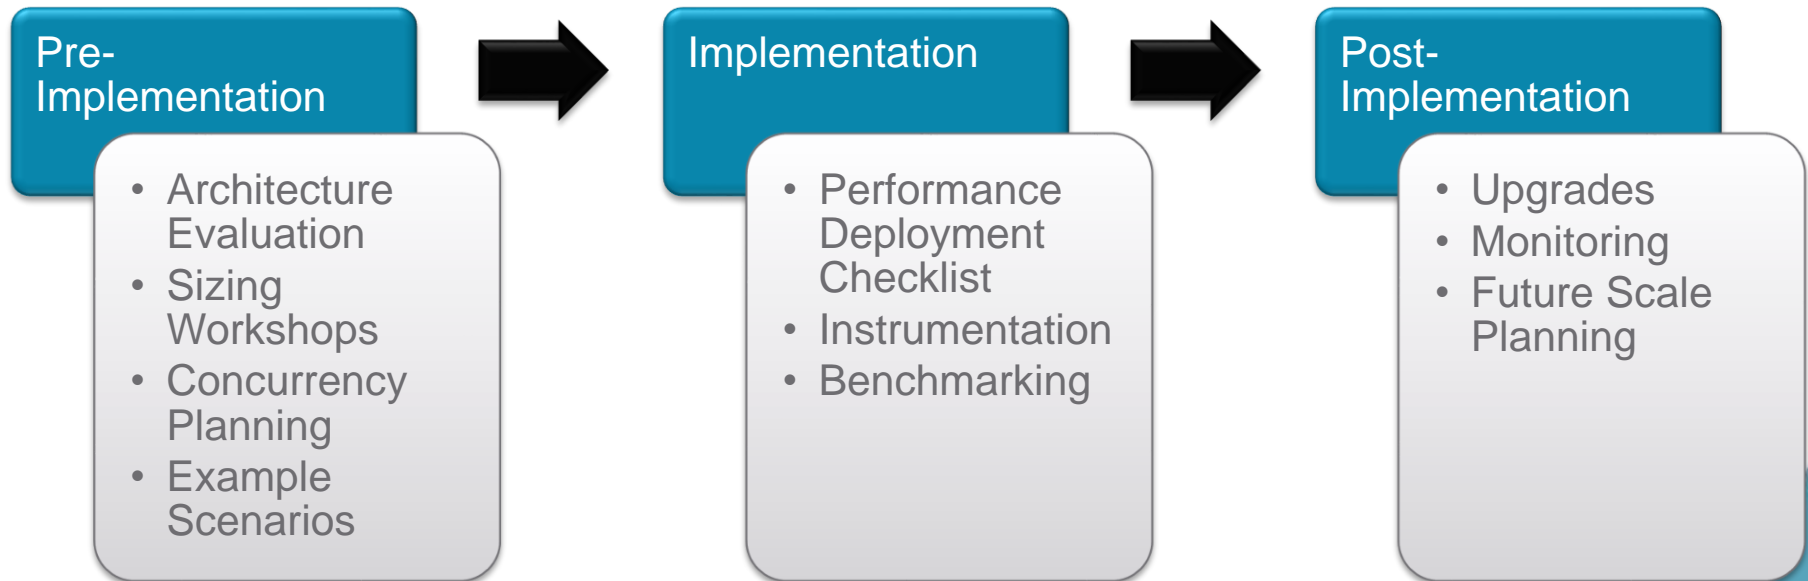
Performance Management

- Performance Management Approach
- Tools & Resources
- Common Pitfalls to Avoid
- Things You Can Do to Improve Performance

Performance Management Approach

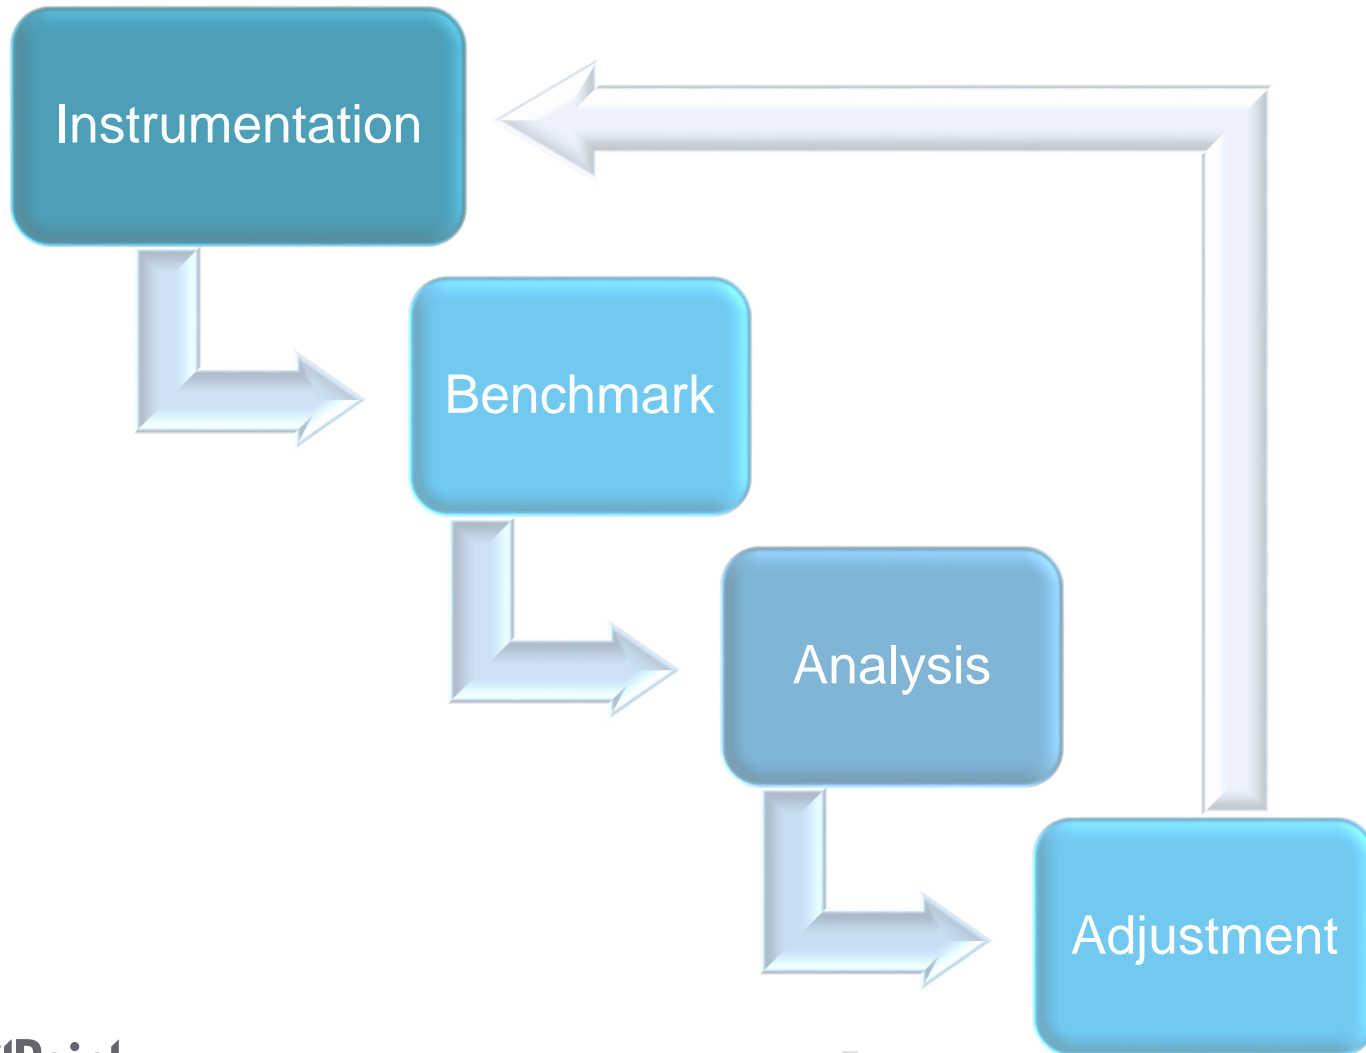
SailPoint Services Approach

- Make performance an integral part of your project
- Don't wait until the end of the project to worry about performance.



Performance Management Methodology

Approach



Resources and Tools

Resources and Tools

Compass – Performance Resource Center

- Performance Optimization Checklist
 - Single sheet review list for staff supporting IdentityIQ
- Performance Management Guide
 - Why and how to change things for the better
- IdentityIQ Hardware Sizing Guide
 - How much hardware is needed to support IdentityIQ?
- Partitioning Best Practices 6.2+
- ...and more!

Resources and Tools

Expert Assistance

- Performance Assessments
 - Architecture analysis
 - Sizing workshops
 - Performance reviews
 - Benchmark and instrument your installation
 - Perform diagnostics on connectors, DB, application servers
 - Provide formal report and recommendations
- Performance Discussion Group
 - Compass → IdentityIQ Forums → Performance

Performance Monitoring Tools

Measuring

- Log4J timings

2015-07-01 09:58:17,314 INFO QuartzScheduler_Worker-4 sailpoint.api.Aggregator:
1481 - Aggregation start

2015-07-01 09:59:24,513 INFO QuartzScheduler_Worker-4 sailpoint.api.Aggregator:
1481 - Aggregation complete

- System out println messages

Performance Monitoring Tools

Measuring Rule Performance

- Meter method

- Track number of calls
- Calculate min/max/avg execution time
- Debug Page → “wrench”
→ Call Timings

- Example code

```
import sailpoint.api.Meter;  
//Begin your code  
Meter.enterByName("MyMeter");  
//Do work  
Meter.exitByName("MyMeter");  
//End of code
```

Call Timings


Name	Hits	Errors	Min (ms)	Max (ms)	Total (ms)	Average (ms)
AccountPolicyExecutor.evaluate	960	0	1	125	1948	2
EntitlementCorrelator.analyzeContributingEntitlements	960	0	0	1	31	0
Identitizer.diff detections	960	0	0	1	8	0
MyMeter	18	0	107	194	2697	149

Performance Monitoring Tools

Task Throughput Report

- Analyzes completed tasks
- Provides benchmark to compare IdentityIQ environments
- Helps diagnose and resolve performance bottlenecks
- Download from Compass

Report Result

Report Details			
Name	Task Througput Details	Started By	spadmin
Type	Live Report	Started	5/13/15 12:27:24 PM
Description	Displays record throughput statistics and performance based on TaskResults.	Completed	5/13/15 12:27:36 PM
Status	 Success	Progress	Download PDF Download CSV

[Return To Reports](#)[Edit Report](#)

Report Data

Task Name	Completed Date	Records Processed	Wall Clock Duration	Aggregate R/S	Working Threads	Partitions Time (ms)	R/S/T	Delete Time	D
Aggregate Accounts - HR	4/8/15 9:11 AM	250025	86.949	2875.536	102	4811607	51.962	0	0
Aggregate Accounts - LDAP	4/8/15 10:26 AM	250016	4451.715	56.161	102	442262665	0.565	0	0

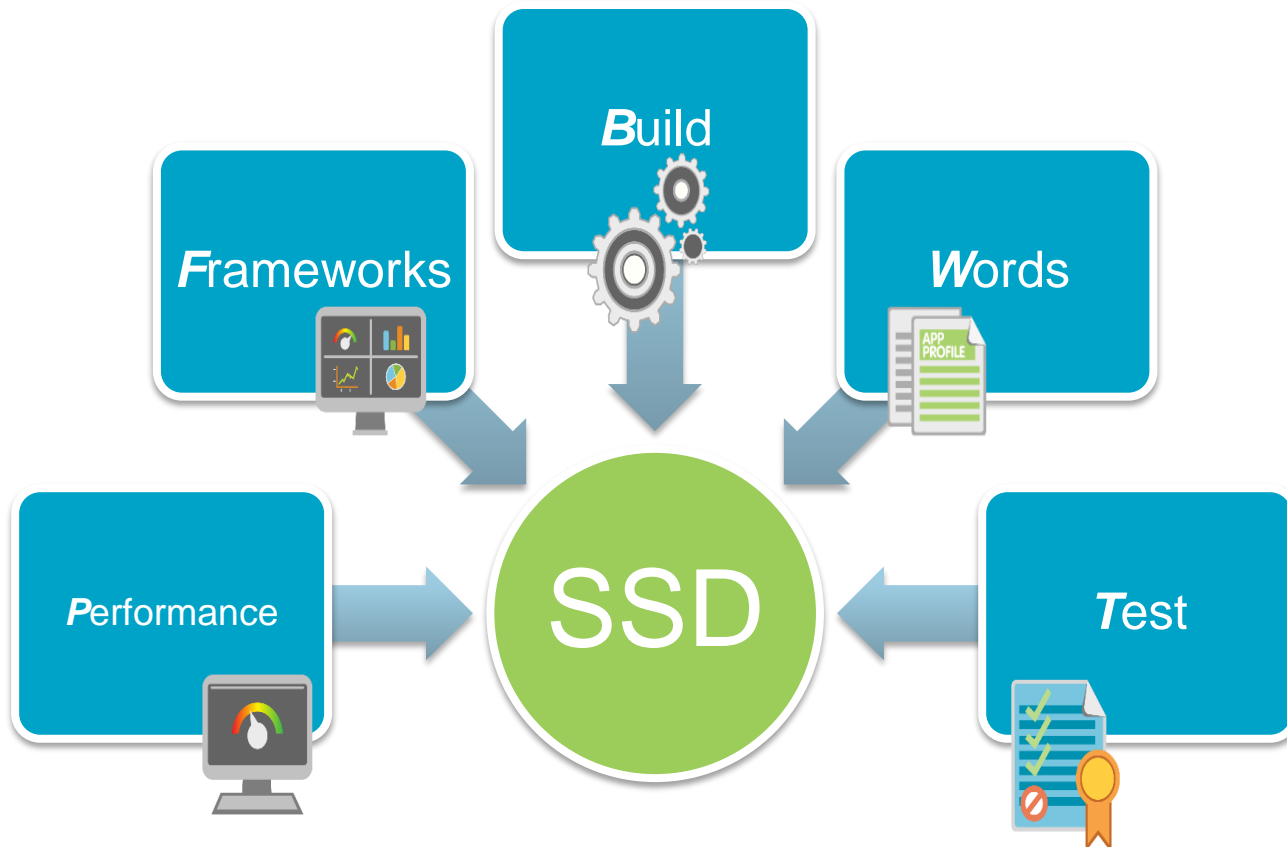
Additional Performance Resources

The Extended Team

- IdentityIQ performance is usually a team effort
- Tuning for optimal performance requires wide expertise
 - Operating systems, databases, networking, Java servlets, J2SE environments, etc.
- Example
 - Database
 - DBA can analyze and tune for performance
 - Application server
 - Admins can monitor threads, connections, etc.
 - Target applications (i.e. LDAP)
 - Admin can analyze and tune for performance

Implementation Tools

Services Standard Deployment (SSD)



Services Standard Deployment

Components

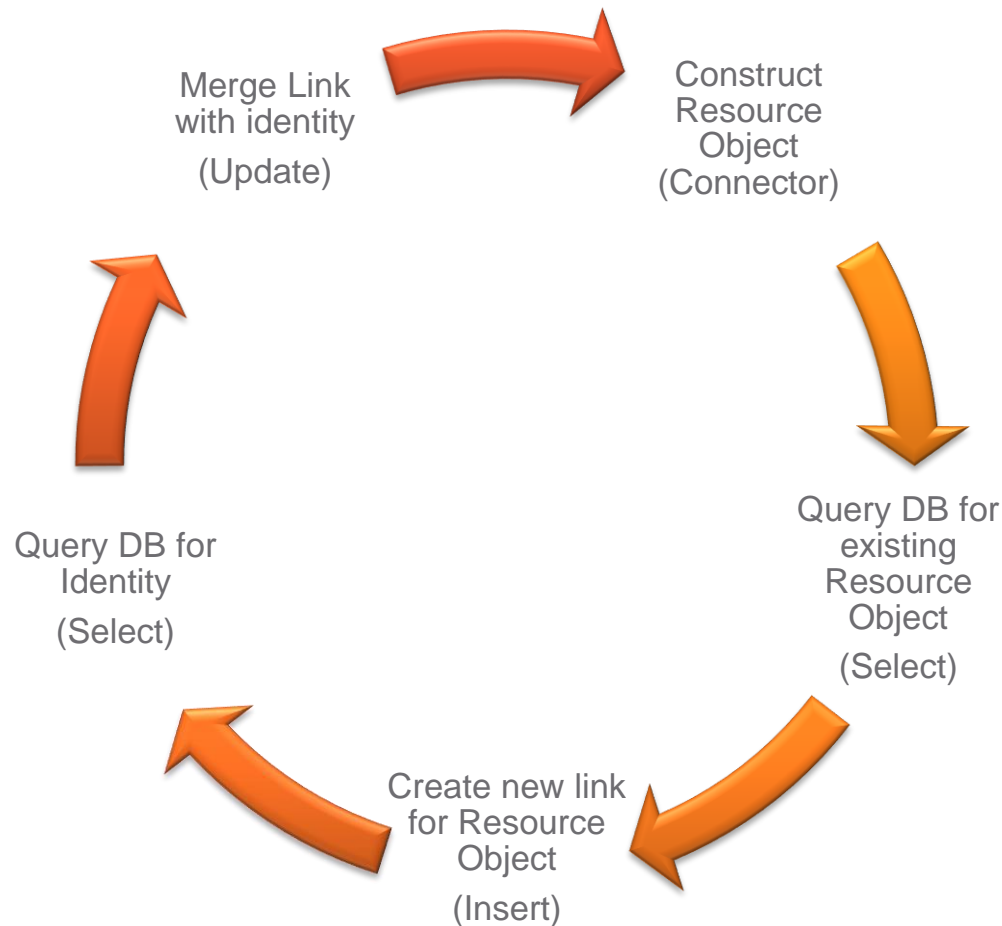
- **Services Standard Build (SSB)**
 - Build site-specific war files, integrate with change control systems
- **Services Standard Framework (SSF)**
 - Collection of workflow templates for extending IdentityIQ to address common provisioning use cases
- **Services Standard Performance (SSP)**
 - Collection of guides and tools used to help measure and modify performance
- **Services Standard Words (SSW)**
 - Document templates for each phase of the project lifecycle
 - Requirements
 - Functional Design
 - Technical Design
 - Unit Testing guide
 - Go Live
- **Available on Compass**

Common Pitfalls

Network, Network, Network

- Traffic between the application and DB server is very sensitive to latency
 - Low Latency connections are best
 - Higher bandwidth connections improve performance
 - Aggregation and certification generations on most network segments should have a .3 ms latency or lower
- Jumbo frames can have a dramatic impact on performance improvement; 1500 MTU is too small for efficient communication in IdentityIQ

Database Example – Aggregation



Virtualization

- Follow VM best practices for applications with large Java heaps
 - Memory reservations
 - CPU reservations
- I/O contention particularly for network resources can represent a severe performance issue
 - Dedicated NICs

Rules/BeanShell

- Lines in rules have a direct correlation to execution time
- Iterative rules have the most impact on performance
 - Move processing out of iterative rules and into non-iterative rules (from Build Map to Pre-Iterate)
 - Use Custom Global or state objects to maintain state during iterations
 - Rules that perform external lookups, joining or other merging of information, will be slow
- Rules that perform correlation or query option searches need to use searchable/indexed attributes
- Rules are more performant than scripts
- No amount of rule tuning can have the impact that responsible rule development can have
- Rule performance profiling is key to a good deployment

Application Server

- Memory, Memory, Memory
- Newer versions of IdentityIQ heavily leverage caching and need memory
- Check hardware against specjbb2005 performance specs, if you aren't getting at least 200K JBO/ps per JVM, performance will be poor
- Tomcat runs with the least amount of overhead, but implement what you know

Database Server


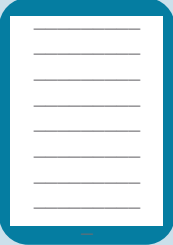





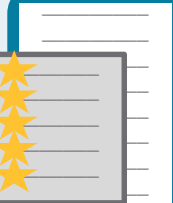
- Watch for I/O contention in shared DB environments
- Ensure that the DB connection pool is adequately sized
 - 50 connections (default)
 - Adequate for sandbox or demos
 - Not typically adequate for production
- Most installations with poor DB performance have typically missed or removed indexes or need additional indexes
- ***Make sure you have the proper and latest JDBC drivers installed (very important!)***
 - See Compass: [JDBC Drivers and IdentityIQ](#)

Application Sources

- Delimited and JDBC datasources should be sorted for optimum data load performance
- Network latency can play a major role in aggregation of JDBC data
 - Consider table replication as an alternative
- LDAP servers may need supplemental indexes to deal with aggregation filters
- API sources (SAP, IDM etc.) are best served via a differential (delta aggregation) or alternative (flatfile/JDBC) approach

Things You Can Do to Improve Performance

Aggregation Strategies

	IdentityIQ	Application
Process All <ul style="list-style-type: none"> Every account read and processed Task option <i>Disable optimization of unchanged accounts = true</i> 		
IdentityIQ-based Optimization (default) <ul style="list-style-type: none"> Every account read Only those with changes are processed Task option <i>Disable optimization of unchanged accounts = false</i> 		
Custom Delta Processing <ul style="list-style-type: none"> Manage own change (i.e. write changed accounts to a flat file and process flat file) Task option <i>Detect deleted accounts = false</i> 		
Connector-based Delta Aggregation* <ul style="list-style-type: none"> Read and process only accounts with changes that have taken place after benchmark <ul style="list-style-type: none"> lastModData, usnChanged, etc. Task option <i>Enable Delta Aggregation = true</i> 		

*Not supported by all connectors

Identity Refresh Strategies

- Identity Refresh task can take time due to processing involved
 - Attribute Refresh
 - Role/Entitlement Processing
 - Policy Violation Checking
 - Risk Calculations
 - Provisioning (attribute sync, role provisioning, process events)
- Refresh only Identities that have changed
 - Delta refresh – aggregation marks cubes with changes for refresh processing
 - Controllable with Task options
 - Filters, time stamp, only update modified identities
- Split Refresh into different tasks
 - Attribute Refresh/Role/Entitlement Processing/Provisioning multiple times a day
 - Policy Checking/Risk Scoring at less frequent intervals
- Divide and Conquer
 - Multi-threading
 - Run on *task* servers
 - Partitioning (mutually exclusive to multi-threading)
 - Run on *request* servers
 - Limit number of threads/partitions – context switching will limit performance

Partitioning


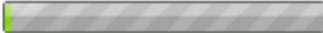
Overview

- Partitioning supports splitting up tasks across multiple servers and threads
- Three items can be partitioned
 - Aggregation
 - All at application level (6.3+)
 - Certain Connectors at connector level (6.2+)
 - LDAP, AD, JDBC, Delimited File, PeopleSoft, SAP
 - Identity Refresh Tasks
 - Manager Certification Generation

Server Object

- View Host Status (read only)
 - “Gear” → Global Settings → Host Configuration

Host Configuration

Hosts					
Host Name	Request Threads	Task Threads	CPU Load	Last Sync	Status
perftest-task1	0	0		12/01/13 10:21:12 pm	● Online
perftest-task2	0	0		12/01/13 10:21:12 pm	● Online

- Heartbeat service updates periodically

Partitioning

View Progress – Task Result

Task Result

Details

Name	Refresh Identity Cube	Started By	The Administrator
Type	Identity	Started	12/1/13 10:36:23 PM
Description	Perform a full refresh of all the identities.	Completed	
Status	pending...	Progress	Beginning identity refresh scan with filter: workgroup == false

[Return to Tasks](#)

Partitioned Results

Name	Host	Status
Refresh 1 to 31,254	perftest-task1	Refreshing 565 1162
Refresh 125,015 to 156,267	perftest-task2	Refreshing 498 94865
Refresh 156,268 to 187,520	perftest-task2	Refreshing 492 133894
Refresh 187,521 to 218,773	perftest-task2	Refreshing 488 110346
Refresh 218,774 to 250,026	perftest-task2	Refreshing 424 149544
Refresh 31,255 to 62,508	perftest-task1	Refreshing 560 8929
Refresh 62,509 to 93,761	perftest-task1	Refreshing 571 16647
Refresh 93,762 to 125,014	perftest-task1	Refreshing 572 24361

Partitioning

Things to be Aware Of

- Limit partitions to fewer than 50
 - Be aware of diminishing returns when scaling up due to overhead
 - context switching, etc.
 - 6.4 aggregation exception: limit partitions to fewer than 250
- Watch for contention
 - DB contention
 - Target application contention
- Ensure ability to terminate tasks
 - For each server, if the Task Processor is running, run a Request Processor to ensure ability to terminate task
- Trial and error testing to fine tune is likely
- More information
 - Compass→IdentityIQ Whitepapers→Partitioning Best Practices

Purge Unneeded Database Objects

- Default for object retention is often “forever”
- These can fill up the database and impact performance
- Set “Days Before...” to non-zero values such as
 - Days before snapshot deletion = 365 (1 yr)
 - Days before task results deletion = 90-180
 - Days before certifications are archived = 720 (2 yrs)
 - Days before certification archive deletion = 1080 (3 yrs)
 - Identity request objects =
 - Task: performidentityrequestmaintenance

Performance Bottom Line

- Build performance into your implementation from the beginning
- Use your resources
 - Performance Team
 - Compass
 - Performance Optimization Checklist
 - Performance Management Guide
 - Sizing Calculator
- Build/leverage your team
 - Pull in resources to tune the infrastructure
- You have a lot of control – develop responsibly

Questions?

Putting it All Together

Section 5, Final Exercise

- Debug and resolve problem discovered during user acceptance testing
- Problem discovered in LCM
- Researching will utilize
 - Identity Cubes
 - Entitlement Catalog
 - Data sources
 - ... and possibly more
- Resolving will utilize
 - Rules
 - Coding
 - Aggregation
 - Refresh
 - ...and possibly more